



## System Call Programming

### Processor Modes

- Mode bit may define areas of memory to be used when the processor is in supervisor mode vs user mode



### System Calls



Trap: System call causes a switch from user mode to kernel mode

#### More examples: System calls

```
struct stat {
    dev_t    st_dev;      /* ID of device containing file */
    ino_t    st_ino;      /* inode number */
    mode_t   st_mode;     /* protection */
    nlink_t  st_nlink;    /* number of hard links */
    uid_t    st_uid;      /* user ID of owner */
    gid_t    st_gid;      /* group ID of owner */
    dev_t    st_rdev;     /* device ID (if special file) */
    off_t    st_size;     /* total size, in bytes */
    blksize_t st_blksize; /* blocksizes for filesystem I/O */
    blkcnt_t st_blocks;   /* number of 512B blocks allocated */

    time_t   st_atime;    /* time of last access */
    time_t   st_mtime;    /* time of last modification */
    time_t   st_ctime;    /* time of last status change */
};
```


## Low Level Process Safety

Questions:

- How do we protect processes from breaking each other? What about breaking the OS?
- Should every process be allowed to execute any command?
- How do we decide what processes deserve which permissions?

### What About User Processes?

- The **kernel** executes privileged operations on behalf of untrusted user processes



### System calls

- A system call involves the following
  - The system call causes a 'trap' that interrupts the execution of the user process (user mode)
  - The kernel takes control of the processor (kernel mode/privilege switch)
  - The kernel executes the system call on behalf of the user process
  - The user process gets back control of the processor (user mode/privilege switch)
- System calls have to be used **with care**.
- Expensive due to **privilege switching**

### Library Functions


- Functions that are a part of standard C library
- To avoid system call overhead use equivalent library functions
  - getchar, putchar vs. read, write (for standard I/O)
  - open, fclose vs. open, close (for file I/O), etc.
- How do these functions perform privileged operations?
  - They make system calls

## Processor Modes

- Mode bit used to distinguish between execution on behalf of OS & behalf of user
- Supervisor mode:** processor executes every instruction in its hardware repertoire
- User mode:** can only use a subset of instructions

### The Kernel


- Code of the OS **executing in supervisor state**
- Trusted software:
  - manages hardware resources (CPU, memory, and I/O)
  - Implements protection mechanisms that could not be changed through actions of untrusted software in user space



### System calls

- `ssize_t read(int fd, void *buf, size_t nbyte)`
  - fd: file descriptor
  - buf: buffer to write to
  - nbyte: number of bytes to read
- `ssize_t write(int fd, const void *buf, size_t nbyte)`
  - fd: file descriptor
  - buf: buffer to write to
  - nbyte: number of bytes to write
- `int open(const char *pathname, int flags, mode_t mode)`
- `int close(int fd)`
- File descriptors:
  - 0 stdin
  - 1 stdout
  - 2 stderr
- Why are these system calls and not just regular library functions?

### So What's the Point?




- Many library functions invoke system calls indirectly
- So why use library calls?
- Usually equivalent library functions make fewer system calls
- non-frequent switches from user mode to kernel mode => less overhead

## Processor Modes

- Instructions can be executed in supervisor mode are supervisor privileges, or protection instruction
  - I/O instructions** are protected. If an application needs to do I/O, it needs to get the OS to do it on its behalf
  - Instructions that can change the **protection state** of the system are privileges (e.g. process' authorization status, pointers to resources, etc)

### The Kernel

**System call interface** is a **safe way** to expose privileged functionality and services of the processor



#### More examples: System calls

- `pid_t getpid(void)`
  - returns the process id of the calling process
- `int dup(int fd)`
  - Duplicates a file descriptor fd. Returns a second file descriptor that points to the same file table entry as fd does.
- `int fstat(int fd, struct stat *buf)`
  - Returns information about the file with the descriptor fd to buf

### Buffering issues

- What is buffering?
- Why do we buffer?
- Can we make our buffer really big?
- What happens if our program exits with a non-empty buffer?

## Unbuffered vs. Buffered I/O

- **Unbuffered**
  - Every byte is read/written by the kernel through a system call
- **Buffered**
  - collect as many bytes as possible (in a buffer) and read more than a single byte (into buffer) at a time and use one system call for a block of bytes

Takeaway: Buffered I/O decreases the number of read/write system calls and the corresponding overhead

## Unbuffered vs. Buffered I/O examples

- **Buffered** output improves I/O performance and can reduce system calls.
- **Unbuffered** output when you want to ensure that the output has been written before continuing.
  - **stderr** under a C runtime library is *unbuffered* by default. Errors are infrequent, but we want to know about them immediately.
  - **stdout** is *buffered* because it's assumed there will be far more data going through it, and more urgent input can use stderr.
  - **logging**: log messages of a process?

## Laboratory

- Write `tr2b` and `tr2u` C programs that transliterate bytes. They take two arguments 'from' and 'to'. The programs will transliterate every byte in 'from' to corresponding byte in 'to'.
  - `./tr2b 'abcd' 'wxyz' < bigfile.txt`
    - Replace 'a' with 'w', 'b' with 'x', etc
  - `./tr2b 'mno' 'pqr' < bigfile.txt`
- `tr2b` uses **getchar** and **putchar** to read from STDIN and write to STDOUT.
- `tr2u` uses **read** and **write** to read and write each byte, instead of using `getchar` and `putchar`. The `nbyte` argument should be 1 so it reads/writes a single byte at a time.
- Test it on a big file with 5000000 bytes

```
$ head --bytes=# /dev/urandom > output.txt
```

## time and strace

- **time** [*options*] *command* [*arguments...*]
- **Output**:
  - `real 0m4.866s`: elapsed time as read from a wall clock
  - `user 0m0.001s`: the CPU time used by your process
  - `sys 0m0.021s`: the CPU time used by the system on behalf of your process
- **strace**: intercepts and prints out system calls to `stderr` or to an output file
  - `$ strace -o strace_output ./tr2b 'AB' 'XY' < input.txt`
  - `$ strace -o strace_output2 ./tr2u 'AB' 'XY' < input.txt`

## Pointers on System Calls

[www.cs.uregina.ca/links/class-info/338/SystemCall\\_IO/SystemCall\\_IO.html](http://www.cs.uregina.ca/links/class-info/338/SystemCall_IO/SystemCall_IO.html)

[courses.engr.illinois.edu/cs241/sp2009/lectures/04-syscalls.pdf](http://courses.engr.illinois.edu/cs241/sp2009/lectures/04-syscalls.pdf)

[www.bottomupcs.com/systen\\_calls.shtml](http://www.bottomupcs.com/systen_calls.shtml)

## Homework 5

- Rewrite `sfrob` using system calls (`sfrobu`)
- `sfrobu` should behave like `sfrob` except:
  - If `stdin` is a regular file, it should initially allocate enough memory to hold all data in the file all at once
  - Consider outputting the number of comparisons performed
- Necessary system calls: `read`, `write`, and `fstat` (read the man pages)
- Measure differences in performance between `sfrob` and `sfrobu` using the **time** command
- Estimate the number of comparisons as a function of the number of input lines provided to `sfrobu`

## Homework 5

- Write a shell script “`sfrobs`” that uses `tr` and the `sort` utility to perform the same overall operation as `sfrob`
- Encrypted input
  - > `tr` (decrypt)
  - > `sort` (sort decrypted text)
  - > `tr` (encrypt)
  - > encrypted output