



Multitasking

- Run multiple processes **simultaneously** to increase performance
- Processes do not share internal structures (stacks, globals, etc)
 - Communicate via **IPC** (inter-process communication) methods
 - Pipes, sockets, signals, message queues
- Single core: Illusion of parallelism** by switching processes quickly (**time-sharing**). **Why is illusion good?**
- Multi-core: True parallelism.** Multiple processes execute **concurrently** on different CPU cores




Multitasking

- `tr -s '[:space:]' '\n' | sort`
- `-u | comm -23 - words`
- Three separate processes spawned simultaneously
 - P1 - tr
 - P2 - sort
 - P3 - comm
- Common buffers (pipes) exist between 2 processes for communication
 - 'tr' writes its stdout to a buffer that is read by 'sort'
 - 'sort' can execute, as and when data is available in the buffer
 - Similarly, a buffer is used for communicating between 'sort' and 'comm'



- A process can be
 - Single-threaded
 - Multi-threaded
- Threads in a process can run in parallel
- A thread is a lightweight process
- It is a basic unit of CPU utilization
- Each thread has its own:
 - Stack
 - Registers
 - Thread ID
- Each thread shares the following with other threads belonging to the same process
 - Code
 - Global Data
 - OS resources (files, I/O)



Single threaded execution

```
int global_counter = 0;
int main()
{
    //code for foo
    ...
    foo(arg1,arg2);
    bar(arg3,arg4,arg5);
    //code for bar
    return 0;
}
```



Multi threaded execution (multiple cores)

```
int global_counter = 0;
int main()
{
    ...
    foo(arg1,arg2);
    bar(arg3,arg4,arg5);
    ...
    return 0;
}
```

```
void foo(arg1,arg2)
{
    //code for foo
}

void bar(arg3,arg4,arg5)
{
    //code for bar
}
```

Multi threaded execution (single core)

```
int global_counter = 0;
int main()
{
    ...
    foo(arg1,arg2);
    bar(arg3,arg4,arg5);
    ...
    return 0;
}
```

```
void foo(arg1,arg2)
{
    //code for foo
}

void bar(arg3,arg4,arg5)
{
    //code for bar
}
```

Time Sharing – Illusion of multithreaded parallelism
(Thread switching has less overhead compared to process switching)

Multithreading properties

- Efficient way to **parallelize** tasks
- Thread switches are less expensive** compared to process switches (context switching)
- Inter-thread communication is easy, via **shared global data**
- Need **synchronization** among threads accessing same data

Pthread API

```
#include <pthread.h>

int pthread_create(pthread_t *thread,
                  const pthread_attr_t *attr, void*
                  (*thread_function) (void*), void *arg);
// Returns 0 on success, otherwise returns non-zero number

void pthread_exit(void *retval);

int pthread_join(pthread_t thread, void **retval);
// Returns 0 on success, otherwise returns non zero error number
```

Thread safety/synchronization

- Thread safe function** - safe to be called by multiple threads at the same time. Function is free of 'race conditions' when called by multiple threads simultaneously.
- Race condition** - the output depends on the order of execution
 - Shared data changed by 2 threads
 - int balance = 1000
 - Thread 1
 - T1 - read balance
 - T1 - Deduct 50 from balance
 - T1 - update balance with new value
 - Thread 2
 - T2 - read balance
 - T2 - add 150 to balance
 - T2 - update balance with new value

Thread safety/synchronization

- Order 1
 - balance = 1000
 - T1 - Read balance (1000)
 - T1 - Deduct 50
 - 950 in temporary result
 - T2 - read balance (1000)
 - T1 - update balance
 - balance is 950 at this point
 - T2 - add 150 to balance
 - 1150 in temporary result
 - T2 - update balance
 - balance is 1150 at this point
 - The final value of balance is 1150
- Order 2
 - balance = 1000
 - T1 - read balance (1000)
 - T2 - read balance (1000)
 - T2 - add 150 to balance
 - 1150 in temporary result
 - T1 - Deduct 50
 - 950 in temporary result
 - T2 - update balance
 - balance is 1150 at this point
 - T1 - update balance
 - balance is 950 at this point
 - The final value of balance is 950

Thread synchronization

- Mutex (mutual exclusion)**
 - Thread 1
 - Mutex.lock()
 - Read balance
 - Deduct 50 from balance
 - Update balance with new value
 - Mutex.unlock()
 - Thread 2
 - Mutex.lock()
 - Read balance
 - Add 150 to balance
 - Update balance with new value
 - Mutex.unlock()
 - balance = 1100
- Only one thread will get the mutex. Other thread will **block in Mutex.lock()**
- Other thread can start execution only when the thread that holds the mutex calls **Mutex.unlock()**

Lab 6

- Evaluate the performance of multithreaded 'sort' command
 - `od -An -f -N 4000000 < /dev/urandom | tr -s ' ' '\n' > random.txt`
 - Might have to modify the command above
- Delete the empty line
 - `time -p sort -g --parallel=2 numbers.txt > /dev/null`
- Add /usr/local/cs/bin to PATH
 - `$ export PATH=/usr/local/cs/bin:$PATH`
- Generate a file containing 10M random **double-precision floating point numbers**, one per line with no white space
 - `/dev/urandom: pseudo-random number generator`

Lab 6

- od**
 - write the contents of its input files to standard output in a user-specified format
 - Options
 - `-t f`: Double-precision floating point
 - `-N <count>`: Format no more than *count* bytes of input
- sed, tr**
 - Remove address, delete spaces, add newlines between each float

Lab 6

- use `time -p` to time the command `sort -g` on the data you generated
- Send output to /dev/null
- Run `sort` with the `--parallel` option and the `-g` option: compare by general numeric value
 - Use `time` command to record the real, user and system time when running `sort` with 1, 2, 4, and 8 threads
 - `$ time -p sort -g file_name > /dev/null (1 thread)`
 - `$ time -p sort -g --parallel={2, 4, or 8} file_name > /dev/null`
 - Record the times and steps in `log.txt`

Ray-Tracing

- Powerful rendering technique in Computer Graphics**
- Yields high quality rendering**
 - Suited for scenes with complex light interactions
 - Visually realistic for a wider variety of materials
 - Trace the path of light in the scene
- Computationally expensive**
 - Not suited for real-time rendering (e.g. games)
 - Suited for rendering high quality pictures (e.g. movies)
- Embarrassingly parallel**
 - Good candidate for **multi-threading**
 - Threads need **not synchronize** with each other, because each thread works on a different pixel (at least at small scale)