

Parallelization

- **Parallelization** is the practice of accelerating a program by running multiple sections simultaneously
- **Process forking** allows for a process to split into multiple subprocesses that run simultaneously
 - Switching between processes (context switching) on the CPU is expensive
 - Inter-process signalling is difficult (eg. pipes)
- **Multithreading** is an efficient type of parallelization
 - **Thread switches are less expensive**
 - Inter-thread signalling is easy via **shared data**
 - Need **synchronization** among threads accessing the same data
 - e.g. Mutex.lock(), Mutex.unlock()

```
#include <pthread.h> ...
const int nthreads = 5;
pthread_t tid[nthreads];
int counter;

void* doSomething(void *arg) {
    counter = counter + (int)arg;
}

int main() {
    int i;
    counter = 0;
    for (i = 1; i <= nthreads; ++i)
        pthread_create(&tid[i], NULL, &doSomething, i);
    for (i = 1; i <= nthreads; ++i)
        pthread_join(tid[i], NULL);
    printf("Counter: %d\n", counter);
    return 0;
}
```

Mutex Example (w/o mutexes)

```
#include <pthread.h>

int pthread_create(pthread_t *thread,
                  const pthread_attr_t *attr, void*
                  (*thread_function) (void*), void *arg);
– Returns 0 on success, otherwise returns non-zero number

void pthread_exit(void *retval);

int pthread_join(pthread_t thread, void **retval);
– thread: thread ID of thread to wait on
– retval: the exit status of the target thread is stored in the
location pointed to by *retval
• Pass in NULL if no status is needed
– Returns 0 on success, otherwise returns non zero error number
```

```
#include <pthread.h> ...
const int nthreads = 5;
pthread_t tid[nthreads];
pthread_mutex_t lock;
int counter;

void* doSomething(void *arg) {
    pthread_mutex_lock(&lock);
    counter = counter + (int)arg;
    pthread_mutex_unlock(&lock);
}
```

Mutex Example (w/ mutexes)

```
int main() {
    int i;
    counter = 0;
    pthread_mutex_init(&lock, NULL);
    for (i = 1; i <= nthreads; ++i)
        pthread_create(&tid[i], NULL, &doSomething, i);
    for (i = 1; i <= nthreads; ++i)
        pthread_join(tid[i], NULL);
    pthread_mutex_destroy(&lock);
    printf("Counter: %d\n", counter);
    return 0;
}
```

Homework 6

- Download the single-threaded raytracer implementation
- Run it to get output image
- Multithread ray tracing
 - Modify main.c and Makefile
- Run the multithreaded version and compare resulting image with single-threaded one

- Build a multi-threaded version of Ray tracer
- Modify "main.c" & "Makefile"
 - Include <pthread.h> in "main.c"
 - Use "pthread_create" & "pthread_join" in "main.c"
 - Link with -lpthread flag (LDLIBS target)
- make clean check
 - Outputs "1-test.ppm"
 - Can see "1-test.ppm"
 - See next slide on how to convert ppm

Ray-Tracing

- **Powerful rendering technique in Computer Graphics**
- **Yields high quality rendering**
 - Suited for scenes with complex light interactions
 - Visually realistic
 - Trace the path of light in the scene
- **Computationally expensive**
 - Not suited for real-time rendering (e.g. games)
 - Suited for rendering high quality pictures (e.g. movies)
- **Embarrassingly parallel**
 - Good candidate for **multi-threading**
 - Threads need **not synchronize** with each other, because each thread works on a different pixel



Simple Example

```
#include <pthread.h> ...
#define NUM_THREADS 5

void *PrintHello(void *thread_num) {
    printf("\nd: Hello World!\n", (int) thread_num);
    pthread_exit(NULL);
}

int main() {
    pthread_t threads[NUM_THREADS];
    int ret, t;
    for(t = 0; t < NUM_THREADS; t++) {
        ret = pthread_create(&threads[t], NULL,
                           PrintHello, (void*)t);
        // check return value
    }
    for(t = 0; t < NUM_THREADS; t++) {
        ret = pthread_join(threads[t], NULL);
        // check return value
    }
}
```

Deadlock

Deadlock:

```
mutex1.lock();      mutex2.lock();
mutex2.lock();      mutex1.lock();
```

What happens if each thread is waiting on a resource that is locked by another?

Solutions

- Ignore (simple to implement, but unsafe)
- Detect (slightly complicated): directed graph cycle checking
- Prevent (very complicated): wait-for-graphs, banker's algorithm, etc.

Race Conditions

Execution order of threads is non-deterministic

Race Condition:

Total = Total + val1	Total = Total - val2
----------------------	----------------------

What value does Total end with?

Solution: Mutexes for synchronization

SIMD vs MIMD

- Multiple Instruction Multiple Data (MIMD)
 - Performs multiple actions on any number of data pieces simultaneously.
 - Standard CPU multithreading (eg. pthread)
- Single Instruction Multiple Data (SIMD)
 - Performs the same action on multiple pieces of data simultaneously.
 - Best for algorithms with little data interaction.
 - Typical of most modern parallel specialized hardware, including GPUs (CUDA).

Ray-tracing





Image Source: POV Ray, Hall of Fame hof.povray.org

Ray-tracing

- Trace the path of a ray from the eye
 - **One ray per pixel** in the view window
 - The color of the ray is the color of the corresponding pixel
- Check for intersection of ray with scene objects.
- **Lighting**
 - Flat shading – The whole object has uniform brightness
 - Lambertian shading – Cosine of angle between surface normal and light direction



Viewing a ppm file

- How to view a ppm file?
 - ppmtojpeg
 - ppmtojpeg - is more lightweight than gimp. If you don't already have it, you can download ppmtojpeg as part of the Netpbm package [here](http://netpbm.sourceforge.net/) (windows,linux,mac)
 - This program comes with many Linux distributions as well as with Cygwin for Windows; it is also installed on the SEAS Unix machines.
 - ppmtojpeg input-file.ppm > output-file.jpg
 - Gimp:
 - sudo apt-get install gimp (Ubuntu)
 - www.gimp.org - or install on your computer (windows,linux,mac)
 - scp the file to your local folder to view it
 - » [conversion tutorial](#) with gimp
 - X forwarding (Inxsv)
 - » gimp 1-test.ppm